



THE AI AGENT BLUEPRINT

How to Build Autonomous AI Agents That Run Your Small
Business

by

Sanjay Shharma

M.Sc. · MBA · APSCM (IIM Calcutta) · 30 years

\$9.99

THE AI AGENT BLUEPRINT

How to Build Autonomous AI Agents That Run
Your Small Business

by

Sanjay Shharma

Director & Promoter, Adorise AI

M.Sc., MBA, APSCM (IIM Calcutta) · 30 years corporate
management

The AI Agent Blueprint

How to Build Autonomous AI Agents That Run Your Small Business

Copyright © 2026 Adorise Digital (a division of Trendy DigiStore LLC).
All rights reserved.

No part of this book may be reproduced, distributed, or transmitted in any form or by any means, including photocopying, recording, or other electronic or mechanical methods, without the prior written permission of the publisher, except in the case of brief quotations embodied in critical reviews and certain other noncommercial uses permitted by copyright law.

AI Disclosure: This book was created with AI-assisted tools for research, drafting, and editing under the direction of the author. All frameworks, examples, and recommendations have been reviewed and approved by the author. The author assumes full responsibility for the accuracy and applicability of the content.

This book is a practical guide based on the author's professional experience. The information contained herein is provided "as is" and without warranty of any kind. The author and publisher disclaim all warranties, express or implied. The reader is responsible for their own outcomes when applying the methods described in this book.

First Edition, 2026

For permissions and inquiries: support@adorisedigital.com

Published by Adorise Digital
30N Gould St, Sheridan, Wyoming, USA
adorisedigital.com

TABLE OF CONTENTS

Chapter 1: PART I: THE FOUNDATION

Chapter 2: CHAPTER 1: WHAT HERMES IS — AND WHY IT CHANGES EVERYTHING

Chapter 3: CHAPTER 2: INSTALLING HERMES — THE 30-MINUTE SETUP

Chapter 4: CHAPTER 3: PROMPTING HERMES — THE AGENT PROMPT FRAMEWORK

Chapter 5: CHAPTER 4: THE CONTENT CREATION WORKFLOW

Chapter 6: CHAPTER 5: THE RESEARCH MONITOR AGENT

Chapter 7: CHAPTER 6: THE CUSTOMER SERVICE AGENT

Chapter 8: CHAPTER 7: THE FINANCIAL REPORTING AGENT

Chapter 9: CHAPTER 8: THE PROPOSAL AND SALES AGENT

Chapter 10: CHAPTER 9: MULTI-AGENT ORCHESTRATION

Chapter 11: CHAPTER 10: MONITORING, OBSERVABILITY, AND MAINTENANCE

Chapter 12: CHAPTER 11: THE 30-DAY IMPLEMENTATION PLAN

HERMES AI AGENT

**The Complete Guide to Building
Autonomous AI Workflows That Run Your
Business While You Sleep**

By Sanjay Shharma

PART I: THE FOUNDATION

CHAPTER 1: WHAT HERMES IS — AND WHY IT CHANGES EVERYTHING

The Day I Stopped Being the Bot

After 30 years in corporate management, I knew the pattern by heart — twelve-hour days, the same emails, the same reports, the same slide decks reformatted for the same meetings. I had spent three decades building a career on doing things well. But somewhere along the way, "doing things well" had become "doing things myself" — even when that meant doing the work a machine could do in seconds.

The breaking point came on a Tuesday at 11:47 PM. I was drafting the 14th version of a quarterly report for a client. The same template. The same data. The same word count. I had been doing this report, with minor variations, every quarter for seven years.

That night I went looking for something — anything — that could break the cycle. A friend mentioned an open-source AI model called Hermes. I downloaded it on a Saturday morning, mostly out of curiosity. I gave it my last quarterly report and asked it to draft a new one in my voice.

It took 90 seconds. The output was 85% there. I edited for 20 minutes. Sent it. The client wrote back: "This is the best one you've done."

That is the day I stopped being the bot. Hermes took over the 60% of my work that was repetitive. I got the 40% back — the strategic thinking, the client relationships, the part of the work that actually requires a human being.

This book is everything I learned in the two years since that night. Every workflow. Every prompt. Every mistake. Every shortcut. The complete guide to building autonomous AI workflows using Hermes — the open-source AI agent that has quietly become the most powerful productivity tool for solo

operators and small teams.

What Exactly Is Hermes?

Hermes is a family of open-source large language models created by Nous Research. Unlike closed models (GPT, Claude, Gemini), Hermes is:

- **Open source** — the model weights are public. You can run it on your own hardware, modify it, fine-tune it, or sell products on top of it
- **Free to use** — no per-token API fees if you self-host
- **Uncensored and unfiltered** — you control the system prompt entirely
- **Functionally capable** — Hermes can call external tools, browse the web, write and execute code, manage files, and chain together multi-step tasks
- **Agentic** — with the right setup, Hermes can complete entire workflows without human intervention

The current generation (Hermes 3, released in 2024, with Hermes 4 in development as of this writing) is competitive with GPT-4 class models on reasoning, coding, and instruction following. The trade-off is speed and infrastructure cost — you need either a powerful local machine, a free cloud tier, or a small paid API budget to run it at production scale.

The key insight that changed my business: Hermes is not a chatbot. Chatbots answer questions. Hermes is an **agent**. Agents complete tasks. That distinction is worth the price of this book.

Chatbot vs. Agent: The Difference That Matters

A chatbot takes a prompt and returns a response. The interaction ends there.

An agent takes a goal and works toward it. The interaction continues until the goal is met or the agent determines it cannot proceed.

Chatbot example:

You: "Write a LinkedIn post about productivity."

Hermes: [returns a single post]

[Interaction over]

Agent example:

You: "Every day at 9 AM, find a trending topic in my industry, write a LinkedIn post about it, generate a 60-second video script, queue the video for rendering, and add the post text to my Buffer queue."

Hermes: [performs the workflow daily, every day, until you turn it off]

[Ongoing autonomous operation]

That second example is not science fiction. It is what this book teaches you to build. By the end, you will have a personal AI agent — running on hardware you control, using a model you can audit, completing real work on your behalf every day.

Who This Book Is For

This book is for:

- **Solo founders and freelancers** who want to scale their output without scaling their hours
- **Small business owners** drowning in repetitive tasks that eat the hours they should be spending on growth
- **Operations managers** looking to automate workflows without hiring a full-time engineer
- **Consultants and coaches** who want to deliver more value to clients without burning out
- **Indie hackers and makers** building SaaS products or content businesses on a budget
- **Curious technologists** who want a practical, project-based introduction to AI agents without the academic overhead

This book is NOT for:

- Machine learning researchers looking for a technical deep-dive into transformer architecture
- Enterprise teams needing SOC2 compliance, audit trails, and dedicated support contracts
- People looking for a one-click "AI employee" solution with no setup (those products exist, but they lock you in)
- Anyone who is not willing to spend 2-3 hours setting up their first workflow

What You Will Build

By the end of this book, you will have built:

A Hermes runtime running on either your local machine or a free cloud service

A multi-model fallback system that switches between providers if one fails

Daily autonomous workflows for content, email, research, and reporting

A content creation pipeline that produces 7+ pieces per week automatically

A research agent that monitors trends, competitors, and news in your industry

A customer service agent that drafts replies for your review

A personal AI dashboard showing you what your agents did and what needs your attention

The total cost to run all of this: \$0-25 per month, depending on your free tier usage.

The time investment to set it up: 10-15 hours spread over 2-3 weekends.

The time you will get back: 10-30 hours per week, starting in week 3.

The Three Philosophies Behind This Book

Philosophy 1: Open source first. You should own the tools you depend on. Closed models can be deprecated, rate-limited, or censored. Hermes is yours. This book shows you how to run it that way.

Philosophy 2: Agents, not toys. AI is only useful when it does real work. Every chapter in this book ends with a working workflow. No theory without implementation. No "you could potentially..." without "here is exactly how to..."

Philosophy 3: Human in the loop. Agents do not replace you. They handle the work you should not be doing, so you can do the work only you can do. Every workflow in this book has a human review step. You stay in control. The agent handles the busywork.

A Note On The 2026 AI Landscape

This book is written in mid-2026. The AI landscape moves fast. New models, new providers, new capabilities appear every month. Here is what has not changed and what has:

Has not changed:

- Hermes is still the leading open-source agent-capable model
- Function calling and tool use are still the foundation of agent workflows
- Prompt engineering is still a real skill
- Cost optimization still matters
- Human review of AI output is still essential

Has changed since 2024:

- Multi-model routing is now standard (use the right model for each task)
- Local inference is fast enough for production on consumer GPUs
- Free cloud tiers are more generous than ever

- The agent framework ecosystem has matured (LangChain, LlamaIndex, AutoGen, CrewAI)
- Voice agents, video agents, and code agents are now mainstream

I will note throughout the book where things are likely to change, and what to do when they do. The workflows are designed to be model-agnostic — if a better model appears next month, you swap it in.

How To Read This Book

The book is structured in three parts:

Part I: The Foundation (Chapters 1-3) — what Hermes is, how to install it, how to write good prompts

Part II: The Workflows (Chapters 4-8) — five real autonomous workflows you can build, with full code

Part III: Scaling & Operations (Chapters 9-11) — multi-agent systems, cost optimization, monitoring, maintenance

If you are impatient, you can skip to Chapter 4 and start building right away. But Chapters 1-3 will save you hours of frustration later. Read them.

If you are methodical, read front to back. Each chapter builds on the previous one.

The appendices contain the full prompt library, tool comparisons, and troubleshooting guides. Keep them open in a separate tab.

Let's get started.

□ KEY TAKEAWAYS

□ This book is written in mid-2026. The AI landscape moves fast. New models, new providers, new capabilities appear every month. Here is what has not changed and what has:

□ **Functionally capable** — Hermes can call external tools, browse the web, write and execute code, manage files, and

chain together multi-step tasks

□ If you are impatient, you can skip to Chapter 4 and start building right away. But Chapters 1-3 will save you hours of frustration later. Read them.

□ **Open source** — the model weights are public. You can run it on your own hardware, modify it, fine-tune it, or sell products on top of it

□ It took 90 seconds. The output was 85% there. I edited for 20 minutes. Sent it. The client wrote back: "This is the best one you've done."

□ Hermes is a family of open-source large language models created by Nous Research. Unlike closed models (GPT, Claude, Gemini), Hermes is:

□ An agent takes a goal and works toward it. The interaction continues until the goal is met or the agent determines it cannot proceed.

CHAPTER 2: INSTALLING HERMES — THE 30-MINUTE SETUP

What You Need

The beauty of Hermes is that you can run it on hardware you already own. No special equipment, no enterprise contracts, no IT department. Here is the complete requirements list:

Minimum (works, but slow):

- Windows 10/11, macOS 12+, or Linux
- 8GB RAM
- 10GB free disk space
- Any CPU from the last 5 years

Recommended (smooth experience):

- 16GB RAM
- 20GB free disk space
- NVIDIA GPU with 6GB+ VRAM (GTX 1660 or better)
- SSD instead of HDD

For local model running:

- NVIDIA GPU with 12GB+ VRAM for Hermes 8B (RTX 3060 or better)
- 24GB+ VRAM for Hermes 70B (RTX 4090 or A6000)

If you do not have a powerful GPU, do not worry. This book shows you how to run Hermes on free cloud tiers that are faster than most local setups. We will cover both paths.

Path A: Local Installation (Ollama — Recommended for Most People)

Ollama is the easiest way to run Hermes on your own machine. It handles model downloads, GPU acceleration, and memory management automatically. The whole setup takes 15 minutes.

Step 1: Install Ollama

On Windows:

Go to <https://ollama.com/download/windows>

Download the installer (about 200MB)

Run the installer — it adds Ollama to your system tray

Open Command Prompt and type: `□`

You should see a version number. If you do, you are good.

On macOS:

Go to <https://ollama.com/download/mac>

Download the .dmg file

Drag Ollama to Applications

Open Terminal and type: `□`

On Linux:

```
curl -fsSL https://ollama.com/install.sh | sh
```

Step 2: Pull the Hermes Model

Now download the actual AI model. This is a one-time download of about 4-8GB depending on which version you choose.

```
ollama pull nousresearch/hermes-2-pro
```

Wait for it to finish. On a 50Mbps connection, it takes about 5 minutes. On a 10Mbps connection, about 25 minutes. Go make coffee.

When it finishes, you will see a success message. You now have a fully functional AI model running locally.

Step 3: Test It


```
ollama run nousresearch/hermes-2-pro "Explain quantum computing in 3 sentences"
```

You should see a streaming response within 2-5 seconds. If you do, congratulations — you have a private AI model running on your own machine. No one is logging your prompts. No rate limits. No monthly fees.

Step 4: Start the API Server

Ollama runs a local API server on port 11434. To enable it:

```
ollama serve
```

This runs in the background. Now any application on your computer can talk to Hermes via:

```
http://localhost:11434/api/generate
```

We will use this in every workflow from Chapter 4 onward.

Path B: Cloud Free Tiers (No Powerful Hardware Needed)

If your computer does not have a GPU, or you want faster responses than local inference can deliver, the cloud free tiers are genuinely excellent in 2026. Here are the top three, ranked:

OpenRouter (Best Overall)

OpenRouter is a unified API that gives you access to dozens of models — including Hermes — through a single endpoint. Free tier models are clearly marked with  in the name.

Setup (3 minutes):

Go to <https://openrouter.ai>

Click "Sign In" → use Google or email

Go to "Keys" → "Create Key" → name it "Hermes-Workflow" → copy the key

Your key looks like: 

Free models available:

- `gpt-4o` — the real Hermes, free
- `gpt-4o-mini` — alternative, fast
- `llama3.1` — good for European languages
- `llama3.3` — newer, very capable

API endpoint:

POST <https://openrouter.ai/api/v1/chat/completions>

Headers:

Authorization: Bearer YOUR_OPENROUTER_KEY

Content-Type: application/json

Body:

```
{  
  "model": "nousresearch/hermes-2-pro:free",  
  "messages": [  
    {"role": "user", "content": "Your prompt here"}  
  ]  
}
```

Rate limits: Around 20 requests per minute on free models. Enough for most personal workflows.

Groq (Fastest Inference)

Groq uses custom LPU chips that are 5-10x faster than GPU inference. Their free tier is generous and very fast.

Setup (3 minutes):

Go to <https://console.groq.com>

Sign up with Google or email

Go to "API Keys" → "Create API Key" → copy

Free tier: 30 requests per minute, 14,400 per day

Models available:

- `gpt-4o` — fast and good
- `gpt-4o-mini` — strong reasoning
- `gemini-1.5-pro` — Google's model, solid

Note: Groq does not host Hermes directly, but the alternatives are comparable for most tasks.

xAI Grok (Generous Free Credits)

xAI gives every new account \$25 in free credits with no credit card required. That is enough for 6+ months of Hermes-level usage.

Setup (5 minutes, requires phone verification):

Go to <https://console.x.ai>

Sign up with X account or email

Verify phone number

\$25 in free credits appears in your dashboard

Go to "API Keys" → "Create New Key" → copy

API endpoint:

```
POST https://api.x.ai/v1/chat/completions
```

Headers:

```
Authorization: Bearer YOUR_XAI_KEY
```

```
Content-Type: application/json
```

Body:

```
{  
  "model": "grok-2-mini",  
  "messages": [  
    {"role": "system", "content": "You are Hermes..."},  
    {"role": "user", "content": "Your prompt"}  
  ]  
}
```

```
}
```

xAI's API is OpenAI-compatible, so any code that works with OpenAI works with Grok.

Path C: Self-Hosted on a Free Cloud VM

If you want the best of both worlds — your own hardware, but better than your laptop — you can spin up a free cloud VM and run Ollama there. This is the most technical path, but it is fully free and gives you a "personal AI server" you can access from anywhere.

Free VM options:

- **Google Colab** — free GPU time, but sessions disconnect after 12 hours
- **Kaggle Notebooks** — 30 hours/week of free GPU
- **Vast.ai** — spot pricing on rented GPUs, sometimes \$0.10-0.30/hour
- **RunPod** — similar to Vast.ai, free tier credits on signup

For most people, this is overkill. Stick with Path A or B unless you have a specific reason to self-host.

My Recommended Setup

For 90% of readers, this is the optimal configuration:

Primary: Local Ollama with Hermes 2 Pro

- Run on your own machine
- No monthly cost
- Private (your prompts never leave your computer)
- Fast enough for personal workflows

Backup: OpenRouter free tier

- Auto-failover if local Ollama goes down
- Use when you need higher quality (the hosted version of Hermes runs on better hardware)

- Use for long-context tasks (100K+ tokens) that local models struggle with

Power user: xAI Grok for batch jobs

- \$25 free credits = 6+ months of heavy use
- Use for high-volume tasks like analyzing 1,000 customer reviews
- Grok is excellent at structured data extraction

We will wire all three together in Chapter 5 with automatic failover logic. If one provider is down or rate-limited, the system automatically tries the next.

Verification: The 5-Minute Test

Whichever path you choose, run this verification before moving on:

Test 1: Basic Generation

```
# Local
```

```
ollama run nousresearch/hermes-2-pro "What is 2+2?"
```

```
# Or via API
```

```
curl -X POST http://localhost:11434/api/generate -d '{
```

```
"model": "nousresearch/hermes-2-pro",
```

```
"prompt": "What is 2+2?",
```

```
"stream": false
```

```
}'
```

Expected: A response containing "4" within 5 seconds.

Test 2: JSON Mode

```
ollama run nousresearch/hermes-2-pro "Output JSON: {\"name\":  
\"Hermes\", \"version\": \"2-pro\"}"
```

Expected: Valid JSON returned.

Test 3: Long Context

```
ollama run nousresearch/hermes-2-pro "Summarize the following
5,000-word article: [paste text]"
```

Expected: A coherent summary within 30 seconds.

If all three pass, you are ready for Chapter 3. If any fail, see the troubleshooting appendix.

Common Setup Mistakes (And How To Fix Them)

Mistake 1: "Model not found" error.

- Cause: You forgot to `pull` the model
- Fix: Run `ollama pull` and wait for completion

Mistake 2: Slow responses (60+ seconds for short prompts).

- Cause: Running on CPU instead of GPU, or model not fully loaded into VRAM
- Fix: Check `ollama ps` to see if model is loaded. Close other heavy applications. If still slow, consider cloud free tier instead.

Mistake 3: "Out of memory" error.

- Cause: Not enough RAM to load the model
- Fix: Use a smaller model variant. Try `qwen2.5:1.5b` (8B) instead of larger versions. Or use cloud free tier.

Mistake 4: Connection refused on port 11434.

- Cause: Ollama API server is not running
- Fix: Run `ollama serve` in a separate terminal. On Windows, it should auto-start. On Linux, you may need to enable the systemd service: `systemctl enable ollama`

Mistake 5: API returns 401 Unauthorized.

- Cause: Wrong API key, or key not passed in the right header

- Fix: Double-check the key. For OpenRouter, the header is `␣`. For xAI, it is `␣`.

What You Have Now

After this chapter, you have:

- `␣` A working Hermes installation (local or cloud)
- `␣` An API endpoint you can call from any application
- `␣` Verified that generation, JSON mode, and long-context all work
- `␣` A backup provider configured (recommended)

In the next chapter, we cover the art of writing prompts that produce consistent, reliable, useful output. This is where most people fail with AI agents — they treat it like a chatbot, not a deterministic system. Chapter 3 fixes that.

The workflows in Chapters 4-8 assume you have completed this setup. Do not skip ahead.

`␣` KEY TAKEAWAYS

- `␣` OpenRouter is a unified API that gives you access to dozens of models — including Hermes — through a single endpoint. Free tier models are clearly marked with ``:free`` in the name.
- `␣` The beauty of Hermes is that you can run it on hardware you already own. No special equipment, no enterprise contracts, no IT department. Here is the complete requirements list:
- `␣` Ollama is the easiest way to run Hermes on your own machine. It handles model downloads, GPU acceleration, and memory management automatically. The whole setup takes 15 minutes.
- `␣` If you do not have a powerful GPU, do not worry. This book shows you how to run Hermes on free cloud tiers that are faster than most local setups. We will cover both paths.
- `␣` We will wire all three together in Chapter 5 with automatic failover logic. If one provider is down or rate-limited, the

system automatically tries the next.

□ Fix: Check □ to see if model is loaded. Close other heavy applications. If still slow, consider cloud free tier instead.

□ Fix: Double-check the key. For OpenRouter, the header is □. For xAI, it is □.

CHAPTER 3: PROMPTING HERMES — THE AGENT PROMPT FRAMEWORK

Why Most People Get Bad AI Output

You have probably seen this pattern. You ask an AI tool for something. The output is generic, hedged, full of disclaimers, and not what you wanted. You try again with a longer prompt. The output is slightly better but still mediocre. You give up and do the task yourself.

The problem is not the AI. The problem is the prompt.

AI models are not mind readers. They default to the most common, most average, most "safe" response unless you specifically instruct them otherwise. That default is fine for casual Q&A. It is useless for production workflows that need consistent, reliable output every time.

After two years of building agent workflows, I have developed a 7-part prompt framework that I use for every autonomous task. It works across every model — Hermes, GPT, Claude, Grok, Llama. Memorize it. Use it. Your AI output will improve 10x.

The 7-Part Agent Prompt Framework

Every prompt you write for an autonomous workflow should contain these 7 sections, in this order:

Role — who the AI is acting as

Context — relevant background information

Task — exactly what to do

Format — how the output should be structured

Constraints — what NOT to do

Examples — sample inputs and outputs

Evaluation — how the AI should check its own work

Let me walk through each section with real examples.

Part 1: Role

The first line of every prompt should assign a specific role. This sets the AI's "voice" and constrains its knowledge base.

Bad:

"Help me write a LinkedIn post."

Good:

"You are an expert B2B SaaS marketing copywriter who specializes in LinkedIn thought leadership for technical founders. You have 10 years of experience writing posts that get 10K+ impressions. Your tone is direct, specific, and never uses corporate jargon."

Why this works: The role section tells the AI not just **what** to do, but **who** would be doing it well. The more specific the role, the more specific the output.

Pro tip: Include a credibility marker ("10 years experience," "published in Harvard Business Review," "specialized in X") to push the AI toward expert-level output. Models are trained to associate these markers with higher quality.

Part 2: Context

Give the AI all the background information it needs to do the task well. This is the section most people skip, and the reason most AI output is generic.

Bad: (no context)

Good:

"Context: I run a 5-person consulting firm that helps mid-size e-commerce companies optimize their customer acquisition cost. My primary audience is founders doing \$1M-10M in annual revenue. I post on LinkedIn 3x per week. My best-performing posts last month were about pricing strategy and customer retention. I avoid AI hype, growth hacking, and

motivational content. I write in first person, use short sentences, and never use emojis."

What to include in Context:

- Your business and what you sell
- Your target audience (specific demographics, not "everyone")
- Your existing content strategy and what has worked
- Your brand voice and tone guidelines
- Any specific examples of past content you are proud of

Pro tip: The more context you give, the less editing you will need to do. Time spent on context saves 10x that time in revisions.

Part 3: Task

Be ruthlessly specific about what you want. Vague tasks produce vague output.

Bad:

"Write a LinkedIn post."

Good:

"Task: Write a LinkedIn post (200-280 words) that:

- Opens with a contrarian take about pricing strategy
- Includes one specific example from a real client engagement (anonymized)
- Makes a non-obvious point that challenges conventional wisdom
- Closes with a specific question that invites comments from founders
- Does NOT use hashtags in the body
- Does NOT end with 'Agree?' or 'Thoughts?'"

Why this works: Every constraint you add removes a degree of freedom from the AI, forcing it toward the specific output you

want. The list format makes the constraints easy to scan and modify.

Pro tip: State negative constraints explicitly. "Do NOT use X" is often more powerful than "Use Y." AI defaults to common patterns. Telling it what to avoid is the fastest way to break those patterns.

Part 4: Format

Tell the AI exactly how to structure the output. This is where you get machine-readable results for downstream automation.

Common formats:

- Plain text (default — usually wrong for workflows)
- JSON (for data extraction, API responses, structured data)
- Markdown (for blog posts, documentation, formatted content)
- CSV (for tabular data, lists, exports)
- HTML (for web content, email templates)
- Specific custom format (e.g., "First line is the title, second line is the subtitle, rest is body")

Example JSON format instruction:

"Output ONLY valid JSON in this exact structure:

```
```
```

```
{
```

```
 \"title\": \"string\",
```

```
 \"body\": \"string\",
```

```
 \"hashtags\": [\"string\"],
```

```
 \"call_to_action\": \"string\"
```

```
}
```

```
```
```

Do not include any text before or after the JSON. Do not wrap it in markdown code blocks. Do not include

any commentary."

Why "ONLY valid JSON" is the most important phrase:

Without it, the AI will often wrap the JSON in explanatory text like "Here is your JSON: ..." which breaks automated parsing.

Part 5: Constraints

This is where you prevent the AI from going off the rails. Constraints are different from format — they are about content, not structure.

Common constraints:

- "Do not make up statistics. If you do not know, say you do not know."
- "Do not use the phrases 'in today's fast-paced world' or 'in the age of.'"
- "Do not apologize. Do not add disclaimers. Do not hedge."
- "Keep the tone professional but human. No corporate jargon."
- "Do not reference being an AI."
- "If the input is unclear, ask for clarification rather than guessing."

Pro tip: Keep a "banned phrase list" that you paste into every prompt. The top offenders:

- "In today's..."
- "It's important to note that..."
- "When it comes to..."
- "At the end of the day..."
- "Leverage" (verb)
- "Synergy"
- "Dive deep" / "Deep dive"
- "Game-changer"
- "Unlock the power of..."
- "In conclusion..."

When you see these in AI output, the AI has defaulted to "average content" mode. Strip them out.

Part 6: Examples

This is the single highest-leverage section of any prompt. Show the AI exactly what good output looks like.

Example for a LinkedIn post prompt:

"Example of the tone and style I want:

>

Post: 'Most pricing pages are lying. I audited 47 SaaS pricing pages last month. 43 of them hid the actual price behind a "Contact Sales" button. The pattern: companies that hide pricing are usually the ones whose pricing is bad. If you cannot say your price on a webpage, you do not know your customer well enough. The fix is not better copywriting. The fix is better positioning. Know your buyer. Set a price that makes sense for them. Show it.'

>

Match this voice: short sentences, specific numbers, contrarian framing, no fluff."

Why one example is not enough: Aim for 2-3 examples. Models pick up on patterns much faster than instructions. "Do X" is weaker than "Here is X done well."

Pro tip: Your examples should be things you actually wrote, not things you wish you wrote. AI learns the **actual** patterns in the examples you give, including subtle things like sentence rhythm and punctuation habits.

Part 7: Evaluation

The final section tells the AI how to check its own work before responding. This is what separates an agent from a chatbot.

Bad: (no evaluation step)

Good:

"Before you output your response, check:

- Does the title contain a number or specific detail? (If not, rewrite it.)
- Is the body under 280 words? (If not, cut it.)
- Is there at least one contrarian or non-obvious point? (If not, add one.)
- Would I be comfortable sending this to my audience without edits? (If not, revise.)

>

Only output your response after you have checked each item."

Why this works: Models are much better at evaluating existing text than generating new text. Asking the AI to check its work against criteria produces output that is 2-3x higher quality on the first try.

Pro tip: Be specific in your checks. "Is it good?" is useless. "Does it contain a number?" is checkable. "Would my audience engage?" is not. Make every check binary and observable.

The Complete Template

Here is the full 7-part template, ready to copy and customize:

ROLE:

You are a [specific role] with [X years] of experience in [specialty].

Your tone is [tone descriptors]. You [what you are known for].

CONTEXT:

[Brief background about your business, audience, brand voice, and

any relevant history.]

TASK:

[Specific, bulleted list of exactly what to do. Include length, structure, must-haves, and must-not-haves.]

FORMAT:

Output ONLY [JSON / Markdown / plain text / CSV] in this structure:

[Exact structure with field names and types]

Do not include any text before or after. Do not wrap in code blocks.

Do not add commentary.

CONSTRAINTS:

- [Banned phrase 1]
- [Banned phrase 2]
- [Behavioral rule 1]
- [Behavioral rule 2]

EXAMPLES:

Here is an example of the style I want:

[Past content you are proud of, with brief annotation explaining

why it works]

EVALUATION:

Before responding, verify:

- [Check 1: binary, observable]
- [Check 2: binary, observable]
- [Check 3: binary, observable]

Only output your response after verifying all checks pass.

Real Example: The Daily Content Workflow Prompt

Here is the actual prompt I use every day to generate my LinkedIn content. It runs autonomously in n8n, gets trends from Google Trends, and produces 3 LinkedIn posts ready for me to review.

ROLE:

You are an expert B2B SaaS marketing copywriter with 10 years of

experience writing LinkedIn thought leadership for technical

founders. You specialize in contrarian takes backed by data.

Your tone is direct, specific, and never uses corporate jargon.

CONTEXT:

I run a 5-person consulting firm that helps mid-size e-commerce

companies optimize their customer acquisition cost. My audience

is founders doing \$1M-10M in annual revenue. I post on LinkedIn

3x per week. My best-performing posts use specific numbers and

challenge conventional wisdom. I write in first person, use

short sentences (avg 12 words), and never use emojis or

inspirational language.

TASK:

Write 3 LinkedIn posts (200-280 words each) on the topic below.

Each post must:

- Open with a contrarian or non-obvious take

- Include at least 2 specific numbers (percentages, dollar amounts, or counts)
- Make one point that challenges conventional wisdom in the industry
- Reference a real-world example or case study
- Close with a specific question (NOT "Thoughts?" or "Agree?")
- Use first person
- No hashtags in body
- No emojis

Topic: {TOPIC_FROM_TRENDS}

Background: {TREND_CONTEXT}

FORMAT:

Output ONLY valid JSON in this exact structure:

```
{  
  "posts": [  
    {  
      "id": 1,  
      "title": "string (compelling, 5-10 words)",  
      "body": "string (200-280 words)",  
      "question": "string (the closing question)",  
      "topic_tags": ["string", "string", "string"]  
    },  
    {  
      "id": 2,  
      "title": "...",
```

```
"body": "...",
"question": "...",
"topic_tags": ["..."]
},
{
"id": 3,
"title": "...",
"body": "...",
"question": "...",
"topic_tags": ["..."]
}
]
}
```

Do not include any text before or after the JSON. Do not wrap in

code blocks. Do not add commentary.

CONSTRAINTS:

- Do not use the phrases: "in today's," "it's important to note,"

"when it comes to," "at the end of the day," "leverage," "synergy,"

"dive deep," "game-changer," "unlock the power of"

- Do not use exclamation marks

- Do not use "you're" or "don't" contractions (use "you are" and

"do not" for slightly more authoritative tone)

- Do not start more than one sentence with "And" or "But"

- Do not apologize or add disclaimers

EXAMPLES:

Here is the style I want to match:

Post 1: "Most pricing pages are lying. I audited 47 SaaS pricing

pages last month. 43 of them hid the actual price behind a 'Contact Sales' button. The pattern: companies that hide pricing

are usually the ones whose pricing is bad. If you cannot say your price on a webpage, you do not know your customer well enough. The fix is not better copywriting. The fix is better positioning. Know your buyer. Set a price that makes sense for them. Show it."

Post 2: "Customer acquisition cost went up 40% last year for e-commerce companies. The standard advice was 'increase your ad spend to maintain reach.' That advice is wrong. CAC is not a function of budget. It is a function of message-audience fit. Increase the second and the first takes care of itself. The companies that maintained profitability in 2025 were the ones that narrowed their targeting, not the ones that broadened it."

Notice: short sentences, specific numbers, contrarian framing, first person, no fluff.

EVALUATION:

Before responding, verify each post:

- Is the word count between 200 and 280?
- Are there at least 2 specific numbers?

- Does it open with a contrarian or non-obvious take?
- Does it close with a specific, non-generic question?
- Are all banned phrases absent?
- Are there any emojis or exclamation marks? (If yes, remove them.)
- Would I send this to my audience without major edits?

Only output your response after verifying all checks pass for all 3 posts.

This prompt is long. That is the point. Long prompts produce reliable, consistent output. Short prompts produce lottery tickets.

The first time you write a prompt like this, it takes 30-45 minutes. The second time, 15 minutes. By the fifth time, you have a template library and you can spin up a new workflow in 5 minutes.

Common Prompting Mistakes (And Fixes)

Mistake 1: Vague role assignment.

Bad: "You are a helpful assistant."

Good: "You are a senior B2B SaaS copywriter with 10 years of experience writing for technical founders. You specialize in contrarian takes backed by data."

Mistake 2: No format specification.

Bad: "Write a product description."

Good: "Output ONLY valid JSON with fields: title (string), description (string, 100-150 words), bullet_points (array of 3 strings), price_suggestion (number)."

Mistake 3: Mixing the task with constraints.

Bad: "Write a post about AI but don't mention ChatGPT and keep it under 200 words and don't sound like a robot."

Good: "Task: Write a post about AI. Constraints: Do not mention ChatGPT. Maximum 200 words. Tone: human, not robotic."

Mistake 4: No examples.

Always include at least one example. The fastest way to improve AI output is to show, not tell.

Mistake 5: Asking the AI to evaluate subjective things.

Bad: "Is this post good?"

Good: "Is the word count between 200-280? Are there at least 2 numbers? Is there a closing question?"

Mistake 6: Forgetting the system prompt vs. user prompt distinction.

The role, context, and constraints go in the **system prompt** (persistent across the conversation). The task, format, examples, and evaluation go in the **user prompt** (specific to this request). Most workflows mess this up and put everything in the user prompt. Use the system prompt for stable instructions.

Quick Reference Card

Save this on your desk or in a note:

PROMPT CHECKLIST:

☐ Role (specific, expert, with credibility)

☐ Context (business, audience, voice, history)

☐ Task (bulleted, specific, with constraints)

☐ Format (exact structure, "ONLY X" rule)

☐ Constraints (banned phrases, behaviors, length)

☐ Examples (1-3 samples, with annotations)

☐ Evaluation (binary checks, observable criteria)

If any section is missing, the output will be worse.

What You Have Now

You have the framework. You can write a prompt that produces consistent, reliable, publication-ready output from any AI model. You have a template you can reuse for every workflow. You have a checklist to make sure no section gets skipped.

In Chapter 4, we start building real workflows. We will use this exact framework to build:

- A daily content creation pipeline
- An email triage agent
- A research monitor
- A customer service reply drafter
- A financial report generator

Each one will be a complete, copy-paste workflow you can run in your own business by the end of the chapter.

□ KEY TAKEAWAYS

□ Good: "Output ONLY valid JSON with fields: title (string), description (string, 100-150 words), bullet_points (array of 3 strings), price_suggestion (number)."

□ Good: "You are a senior B2B SaaS copywriter with 10 years of experience writing for technical founders. You specialize in contrarian takes backed by data."

□ Give the AI all the background information it needs to do the task well. This is the section most people skip, and the reason most AI output is generic.

□ This is where you prevent the AI from going off the rails. Constraints are different from format — they are about content, not structure.

□ This prompt is long. That is the point. Long prompts produce reliable, consistent output. Short prompts produce lottery tickets.

□ The first line of every prompt should assign a specific role. This sets the AI's "voice" and constrains its knowledge base.

□ The final section tells the AI how to check its own work before responding. This is what separates an agent from a

chatbot.

ABOUT THE AUTHOR & THE AI MASTERY FLEET

Executive Leadership, Autonomous Architecture & Enterprise B2B Systems

ENJOYED THIS FREE OPENING SAMPLE?

You have just experienced the foundational chapters of **The AI Agent Blueprint**. To unlock the complete, unabridged masterwork (including all production code snippets, architecture schemas, system prompts, and client onboarding SOPs), order instant digital access or print editions below:

- **Official Digital Store:** <https://books.adorisedigital.com> (Instant PDF + Kindle EPUB)
- **Global Digital Storefront:** <https://adorisedigital.gumroad.com> (Worldwide Instant Access)
- **Paperback & Kindle on Amazon:** Search Google for "*The AI Agent Blueprint Sanjay Shharma*"

Sanjay Shharma is an Author, Corporate Executive Mentor, and Technology Strategist holding an M.Sc., MBA, and APSCM from the prestigious **Indian Institute of Management (IIM) Calcutta**, with over thirty years of executive management and business leadership.

He architected **The AI Mastery & Receptionist Series** because modern enterprises are overwhelmed by repetitive operational bottlenecks, missed high-ticket phone calls, and high employee churn. Instead of theoretical AI jargon, his frameworks provide battle-tested, zero-code and low-code automated software fleets that capture revenue 24/7.

Author Credentials & Strategic Ecosystem:

Academic Alumni	Indian Institute of Management Calcutta (IIM Calcutta) • APSCM
Specialization	Autonomous AI Agents, Voice AI Receptionists & B2B Operational Scaling
Publisher	Adorise Digital • A Division of Trendy DigiStore LLC (Wyoming, USA)
Official Bookstore	https://books.adorisedigital.com
Global Storefront	https://adorisedigital.gumroad.com
Author LinkedIn	linkedin.com/in/sanjayshharma
Enterprise Portal	https://adorisedigital.com
B2B & Reader Support	support@adorisedigital.com

TURNKEY ENTERPRISE AI IMPLEMENTATION:

Need turnkey autonomous voice receptionists, lead capture bots, or programmatic marketing pipelines deployed for your firm? Explore live software architectures and book an executive consultation at adorisedigital.com or email support@adorisedigital.com.

THE 48-BOOK MASTER AUTHOR CATALOG • PART 1

Flagship Works by Sanjay Shharma • Published Worldwide by Adorise Digital

1. THE CARE COLLECTION (14 Flagship Volumes • Marital Health, Aging & Healing)

- #1 **Still Married, Just Silent** — 14 conversations to reopen an emotionally quiet marriage
- #2 **The Quiet Care** — How to give and receive care in a marriage that lasts
- #3 **The Second Chapter** — Rebuilding identity, purpose, and daily joy after caregiving
- #4 **The Second Honeymoon** — Reigniting romance, vulnerability, and play after kids leave
- #5 **The Survivor** — Compassionate protocol when one spouse is left behind
- #6 **The Long Hello** — How couples 50–70 reconnect after the last child leaves
- #7 **When They Forget Your Name** — Navigating dementia, cognitive decline, and memory care
- #8 **Gray Divorce** — Coming back together (or parting) with financial dignity at 55+
- #9 **The Last Sandwich** — Caring for aging parents, spouse, and self at the same time
- #10 **The Long Goodbye** — When your spouse is dying and the marriage is still yours
- #11 **Remarried at 60** — Blending adult families, estates, and hearts in sunset years
- #12 **After the Last Goodbye** — Practical and emotional roadmap for first 365 days of loss
- #13 **The Friends Who Fade** — Why friendships disappear in midlife & building adult tribes
- #14 **Hum Ladte Nahi (Hinglish)** — Shaadi ko wapas kholne ka 30-din ka practical plan (India)

2. THE AI MASTERY & VOICE AUTOMATION FLEET (10 Volumes • Agency & Operations)

- #15 **The AI Agent Blueprint** — Building, deploying, and monetizing autonomous software fleets
- #16 **The AI Agency Playbook** — From zero to \$50,000/month offering turnkey AI automation
- #17 **The AI Content Agent** — Autonomous video, social scheduling, and programmatic SEO
- #18 **The AI Agent Action Workbook** — Interactive architecture schemas, code snippets, and worksheets
- #19 **The AI Agent Cheat Sheets** — 50+ battle-tested system prompts for client deployment
- #20 **Never Miss Another Call** — How front-desk voice AI captures \$100,000 in lost revenue
- #21 **The AI Receptionist Blueprint** — Architecture, Twilio SIP trunking, and CRM integration
- #22 **The AI Receptionist Buyer's Guide** — What founders must know before buying or deploying
- #23 **AI Front Desk for HVAC & Home Services** — Capturing emergency dispatch calls 24/7 without missed leads
- #24 **AI Front Desk for Law Firms & Clinics** — Qualifying high-net-worth client consultations automatically

Catalog Continues on Next Page: Indian Wealth Series & Official Global Buying Links →

THE 48-BOOK MASTER AUTHOR CATALOG • PART 2

Flagship Works by Sanjay Shharma • Published Worldwide by Adorise Digital

3. THE INDIAN WEALTH & FINANCIAL MASTERY SERIES (24 Volumes • 8 Frameworks)

Books 25-27	CIBIL Score Repair Blueprint: Vol 1: 7-Step Audit Vol 2: Tactical Manual Vol 3: Legal Notices
Books 28-30	Home Loan Approval Playbook: Vol 1: Eligibility Audit Vol 2: Tactical Manual Vol 3: Banking Dossier
Books 31-33	Business Loan / MSME Finance: Vol 1: CGTMSE Audit Vol 2: Tactical Manual Vol 3: DSCR Templates
Books 34-36	Credit Card Optimization: Vol 1: APR & Fee Audit Vol 2: Tactical Manual Vol 3: Dispute Kit
Books 37-39	Debt Settlement & Relief: Vol 1: OTS Feasibility Vol 2: Tactical Manual Vol 3: Legal Defense
Books 40-42	Global Remote Career (USD): Vol 1: Resume & ATS Vol 2: Tactical Manual Vol 3: Tax & Wire Setup
Books 43-45	GST Filing & Reconciliation: Vol 1: GSTR-1/3B Audit Vol 2: Tactical Manual Vol 3: Notice Defense
Books 46-48	GST Export & Inverted Refunds: Vol 1: RFD-01 Audit Vol 2: Tactical Manual Vol 3: Appellate Appeals

WHERE TO GET THE BOOKS • OFFICIAL BUYING LINKS

1. Official Digital Bookstore (Direct DRM-Free Delivery):

- Instant access to Master PDFs, Kindle EPUBs, Audiobooks & Action Toolkits.
- Domestic Indian Payments (UPI / QR / NetBanking): <https://books.adorisedigital.com>
- Global Storefront (Cards / Apple Pay / PayPal / Whop): <https://adorisedigital.gumroad.com>

2. Amazon KDP & Google Play Books:

- Paperback, Hardcover, and Kindle editions available on Amazon worldwide.
- Simply search Google for: “[Book Title] Sanjay Shharma” or “The AI Agent Blueprint Sanjay Shharma”

3. Direct Publisher Support & Author Ecosystem:

- Email Support: support@adorisedigital.com | Portal: adorisedigital.com
- Published by **Adorise Digital** • A Division of Trendy DigiStore LLC, Sheridan, Wyoming, USA