



THE AI CONTENT AGENT

Daily Posts, Videos, and Newsletters on Autopilot

by

Sanjay Shharma

M.Sc. · MBA · APSCM (IIM Calcutta) · 30 years

\$9.99

THE AI CONTENT AGENT

Daily Posts, Videos, and Newsletters on Autopilot

by

Sanjay Shharma

Director & Promoter, Adorise AI

M.Sc., MBA, APSCM (IIM Calcutta) · 30 years corporate
management

The AI Content Agent

Daily Posts, Videos, and Newsletters on Autopilot

Copyright © 2026 Adorise Digital (a division of Trendy DigiStore LLC).
All rights reserved.

No part of this book may be reproduced, distributed, or transmitted in any form or by any means, including photocopying, recording, or other electronic or mechanical methods, without the prior written permission of the publisher, except in the case of brief quotations embodied in critical reviews and certain other noncommercial uses permitted by copyright law.

AI Disclosure: This book was created with AI-assisted tools for research, drafting, and editing under the direction of the author. All frameworks, examples, and recommendations have been reviewed and approved by the author. The author assumes full responsibility for the accuracy and applicability of the content.

This book is a practical guide based on the author's professional experience. The information contained herein is provided "as is" and without warranty of any kind. The author and publisher disclaim all warranties, express or implied. The reader is responsible for their own outcomes when applying the methods described in this book.

First Edition, 2026

For permissions and inquiries: support@adorisedigital.com

Published by Adorise Digital
30N Gould St, Sheridan, Wyoming, USA
adorisedigital.com

TABLE OF CONTENTS

Chapter 1: PART 1: THE CONTENT CREATION PROBLEM

Chapter 2: Chapter 1: Why Most Creators Burn Out (And How AI Changes It)

Chapter 3: Chapter 2: The Content Engine Architecture

Chapter 4: PART 2: THE PRACTICAL WORKFLOWS

Chapter 5: Chapter 3: The 2-Hour Sunday Batch Session

Chapter 6: Chapter 4: Platform-Specific Workflows

Chapter 7: PART 3: THE AI STACK

Chapter 8: Chapter 5: Setting Up Hermes for Content

Chapter 9: Chapter 6: The Content Prompts (The Real Engine)

Chapter 10: PART 4: THE PUBLISHING SYSTEM

Chapter 11: Chapter 7: The Scheduling Stack

Chapter 12: Chapter 8: The Analytics Loop

Chapter 13: PART 5: THE LONG-TERM GAME

Chapter 14: Chapter 9: The 12-Month Content Roadmap

Chapter 15: Chapter 10: Common Mistakes (And How to Avoid Them)

Chapter 16: Chapter 11: The Hermes Content Stack (Reference)

Chapter 17: PART 6: ADVANCED CONTENT STRATEGY

Chapter 18: Chapter 12: The Advanced Content Strategy (For 100K+ Audience)

Chapter 19: PART 7: 12-MONTH CONTENT ROADMAP

Chapter 20: Chapter 13: The 12-Month Content Roadmap

Chapter 21: PART 8: REAL-WORLD CASE STUDIES

Chapter 22: Chapter 14: 4 Content Creator Case Studies

Chapter 23: CHAPTER 12: THE 15 CONTENT CASE STUDIES

Chapter 24: CHAPTER 13: THE 20 CONTENT AGENT MISTAKES

Chapter 25: CHAPTER 14: THE 12-MONTH CONTENT ROADMAP

HERMES AI FOR CONTENT CREATORS

**Daily Posts, Videos, and Newsletters on
Autopilot**

PART 1: THE CONTENT CREATION PROBLEM

Chapter 1: Why Most Creators Burn Out (And How AI Changes It)

The average content creator publishes 3-5 pieces of content per week. The successful creator publishes 7-15. The prolific creator publishes 20-30. The math is brutal. At 15 minutes per piece, 15 pieces per week is 3.75 hours per week. At 30 minutes per piece, 15 pieces is 7.5 hours per week. The successful creator is spending 5-10 hours per week on content creation alone, on top of the work of actually running the business.

This is why most creators burn out. The content treadmill is relentless. The audience wants daily posts. The algorithms reward consistency. The competition is publishing every day. The creator who takes a week off loses 10% of their reach. The creator who takes a month off loses 50%.

This book solves the content treadmill with AI. Specifically, with Hermes — the open-source AI model from Nous Research. Hermes is fast, runs locally or on cheap GPUs, and is fully customizable. With Hermes as your content engine, you can publish 15-30 pieces per week while spending 3-5 hours per week on the actual creative work.

This is not a book about replacing yourself with AI. This is a book about using AI to amplify your voice. The creative direction, the strategy, the brand voice — those are still yours. The execution, the formatting, the repurposing — those are the AI's job. You focus on what makes the content yours. The AI handles what makes the content consistent.

The Content Treadmill

The content treadmill has 3 components:

Ideation — Coming up with new content ideas

Creation — Writing, recording, designing the content

Distribution — Publishing across platforms, repurposing for different formats

Most creators spend their time on creation. The math: 60% of time on creation, 20% on ideation, 20% on distribution. With AI, the math changes to: 20% on creation, 30% on ideation (because AI generates ideas faster than you can think), 50% on distribution (because AI enables 5-10x more distribution through repurposing).

The result is the same content output with 30-50% less time. The creator who spends 10 hours per week on content can now spend 5-7 hours per week and produce 2-3x more content.

What Hermes Does Well (And What It Doesn't)

Hermes is a large language model. It is not a search engine. It is not a real-time data source. It is not a content idea generator for breaking news. Hermes is excellent at:

- **Repurposing** — Turn a long article into 10 social posts, 5 emails, 3 video scripts
- **Drafting** — Write a first draft of a blog post, video script, or social post
- **Summarizing** — Condense a 5,000-word article into a 280-character tweet
- **Formatting** — Convert bullet points to prose, prose to bullets, etc.
- **Translating** — Convert content to 10+ languages for global reach
- **Variations** — Generate 5 variations of a headline, opening, or hook
- **Editing** — Proofread, tighten, and improve existing content

Hermes is not great at:

- Real-time information (use Google, not Hermes)
- Subject matter expertise (use a domain expert or your own knowledge)

- Original research (use actual research, not LLM guesses)
- Personal voice (your voice comes from editing the AI's output, not from raw AI generation)

The right mental model: Hermes is a brilliant intern who can draft 80% of any content in 30 seconds. Your job is the 20% that makes it yours. The 80/20 split is the leverage.

The Hermes Stack

The Hermes stack for content creators has 4 layers:

Layer 1: Hermes base model. The open-source model from Nous Research. Available in 7B, 13B, and 70B parameter versions. Runs on local GPUs (7B/13B) or cloud GPUs (70B). Cost: free (local) or \$0.50-2.00/hour (cloud).

Layer 2: Open Hermes fine-tunes. Domain-specific fine-tunes (coding, math, chat). The Open Hermes 2.5 model is the recommended starting point for content work.

Layer 3: Orchestration tools. Tools that chain multiple LLM calls. Examples: LangChain, LlamaIndex, Haystack. These tools let you build content pipelines (input → research → draft → edit → format).

Layer 4: Publishing tools. The final step — publishing to your platforms. Examples: Buffer, Hypefury, Zapier. These tools take AI-generated content and publish to social, email, blog, etc.

The stack is not magic. The stack is a system. The system requires setup. Once set up, the system runs with minimal intervention.

The 3 Content Modes

There are 3 content modes creators operate in:

Mode 1: Reactive (most creators). Create content based on what feels urgent. React to trends, news, comments. This mode produces inconsistent output and burns out the creator.

Mode 2: Scheduled (better creators). Create content on a fixed schedule. Monday: blog post. Tuesday: video. Wednesday:

email. This mode produces consistent output but requires 5-10 hours per week.

Mode 3: Systematic (AI-powered creators). Create content in batches using AI. Sunday: 2-hour batch session produces 15-30 pieces for the week. AI handles the variations, the repurposing, the formatting. This mode produces the most output with the least time.

This book focuses on Mode 3. The systematic creator uses AI to batch-create content. The batch session is the only time the creator needs to be "on" creatively. The rest of the week, the AI handles the execution.

□ KEY TAKEAWAYS

- The result is the same content output with 30-50% less time. The creator who spends 10 hours per week on content can now spend 5-7 hours per week and produce 2-3x more content.
- The right mental model: Hermes is a brilliant intern who can draft 80% of any content in 30 seconds. Your job is the 20% that makes it yours. The 80/20 split is the leverage.
- Hermes is a large language model. It is not a search engine. It is not a real-time data source. It is not a content idea generator for breaking news. Hermes is excellent at:
- The stack is not magic. The stack is a system. The system requires setup. Once set up, the system runs with minimal intervention.
- Personal voice (your voice comes from editing the AI's output, not from raw AI generation)
- **Repurposing** — Turn a long article into 10 social posts, 5 emails, 3 video scripts
- **Distribution** — Publishing across platforms, repurposing for different formats

Chapter 2: The Content Engine Architecture

The content engine is the system that produces your content week after week. This chapter covers the architecture.

The 4-Stage Pipeline

The content engine has 4 stages:

Stage 1: Idea Generation (Sunday, 30 min). The creator generates 30-50 content ideas for the week. Ideas come from: customer questions, industry trends, evergreen topics, content pillars. AI generates variations and angles for each idea.

Stage 2: Content Creation (Sunday, 60-90 min). The creator turns 5-10 of the best ideas into first drafts. AI handles the bulk of the writing. The creator edits for voice, accuracy, and brand alignment.

Stage 3: Repurposing (Monday, 30-60 min). The creator takes the 5-10 first drafts and uses AI to repurpose each one into 5-10 platform-specific pieces. One blog post becomes 10 tweets, 3 LinkedIn posts, 2 email newsletters, 1 video script, 5 Instagram captions.

Stage 4: Scheduling and Publishing (Monday, 30 min). The creator uses a scheduling tool (Buffer, Hypefury) to queue the 50-100 pieces for the week. The publishing happens automatically over the next 7 days.

Total creator time: 2.5-3.5 hours per week. Total content output: 50-100 pieces per week.

The Content Pillars

Before the engine runs, the creator needs content pillars. Content pillars are the 3-5 core topics the creator publishes about. Every piece of content fits into a pillar.

Example content pillars for a marketing consultant:

- Pillar 1: Lead generation strategies
- Pillar 2: Conversion optimization
- Pillar 3: Marketing automation
- Pillar 4: Client retention
- Pillar 5: Agency building

Every blog post, tweet, video, or email fits into one of these 5 pillars. The pillars are the brand. The pillars are the content strategy.

The pillar framework:

List 5 topics you can talk about for 30 minutes without notes

Each topic becomes a pillar

Each pillar gets 20% of the content calendar

AI generates variations within each pillar

The result is consistent, focused content

The Content Calendar

The content calendar is the publishing schedule. The calendar has 3 layers:

Layer 1: Anchor content (1-2 per week). Long-form content: blog post, video, podcast episode. This is the high-effort, high-value content. The creator writes or records this personally.

Layer 2: Derivative content (5-10 per week). Pieces derived from the anchor content. One blog post becomes 5-10 social posts, 2-3 emails, 1 video script. The AI does most of the work. The creator edits.

Layer 3: Reactive content (5-10 per week). Real-time content: replies to comments, engagement on other creators' posts, news reactions. The creator does this manually, but with AI assistance for drafting.

Total: 15-30 pieces per week. Creator time: 2.5-3.5 hours. AI handles 80% of the work.

The AI Prompts (The Most Important Part)

The content engine runs on prompts. The 5 core prompts are:

Prompt 1: Idea generation. "Give me 30 content ideas for a [niche] creator targeting [audience]. Each idea should be specific, actionable, and fit a [pillar]."

Prompt 2: Draft writing. "Write a [length] [format] about [topic] for [audience]. The tone should be [tone]. The hook should be [hook]. Include [specific elements]."

Prompt 3: Repurposing. "Take this [format] about [topic] and turn it into [list of target formats]. Keep the core message. Adjust the tone for each platform."

Prompt 4: Headline variations. "Give me 10 variations of this headline: [headline]. Each variation should be [length] words and use a different psychological hook: curiosity, urgency, specificity, controversy, benefit, fear, social proof, authority, scarcity, transformation."

Prompt 5: Engagement. "Reply to this comment in a way that adds value, builds relationship, and is consistent with the creator's voice: [comment]."

The prompts are the engine. The prompts are trained on examples of your best content. The prompts are refined based on what works.

□ KEY TAKEAWAYS

□ Before the engine runs, the creator needs content pillars. Content pillars are the 3-5 core topics the creator publishes about. Every piece of content fits into a pillar.

□ Every blog post, tweet, video, or email fits into one of these 5 pillars. The pillars are the brand. The pillars are the content strategy.

□ The prompts are the engine. The prompts are trained on examples of your best content. The prompts are refined based on what works.

- The content engine is the system that produces your content week after week. This chapter covers the architecture.
- Total creator time: 2.5-3.5 hours per week. Total content output: 50-100 pieces per week.
- Total: 15-30 pieces per week. Creator time: 2.5-3.5 hours. AI handles 80% of the work.
- The content calendar is the publishing schedule. The calendar has 3 layers:

PART 2: THE PRACTICAL WORKFLOWS

Chapter 3: The 2-Hour Sunday Batch Session

This chapter walks through the actual 2-hour Sunday batch session. The session produces 50-100 pieces of content for the week. The session has 4 phases.

Phase 1: Idea Generation (20 minutes)

Setup: Open your AI tool (Hermes local or via API). Have your content pillars in front of you.

Process:

For each pillar, prompt the AI: "Give me 20 content ideas for [pillar] targeting [audience]. Each idea should be specific, not generic."

Review the 100 ideas (5 pillars × 20 ideas). Mark the top 30.

Group the 30 ideas by format: blog posts, videos, social posts, emails.

Final selection: pick 5-10 ideas for anchor content + 20-30 ideas for derivative content.

Output: A spreadsheet with 30-50 ideas, grouped by format, ready for drafting.

Phase 2: Anchor Content Creation (60 minutes)

Setup: Open your writing tool (Google Docs, Notion). Have the AI prompt ready.

Process:

Pick the highest-priority anchor content (1-2 pieces).

For each piece, prompt the AI to generate a first draft.

Edit the draft for voice, accuracy, brand alignment (20-30 minutes per piece).

Format for publishing: blog, video script, podcast outline.

Add images, links, and calls-to-action.

Output: 1-2 anchor pieces (blog post, video, podcast) ready for the week.

Phase 3: Repurposing (30 minutes)

Setup: Open your repurposing tool. Have the 1-2 anchor pieces ready.

Process:

For each anchor piece, prompt the AI to repurpose into 5-10 platform-specific pieces:

- 10 tweets (Twitter/X)
- 3 LinkedIn posts (longer form)
- 2 email newsletter versions
- 1 short-form video script (Reels, TikTok, Shorts)
- 5 Instagram captions
- 3-5 carousel slides (LinkedIn, Instagram)

Edit each piece for platform-specific voice and format.

Save to scheduling tool or content library.

Output: 30-50 derivative pieces ready for scheduling.

Phase 4: Scheduling (10 minutes)

Setup: Open your scheduling tool (Buffer, Hypefury, Later).

Process:

Upload the 30-50 pieces to the scheduling tool.

Set the publishing schedule: 3-5 pieces per day across the week.

Review the queue for balance: enough variety, not too much repetition, good distribution across pillars.

Hit "Schedule."

Output: 30-50 pieces scheduled to publish over 7 days. Total time: 2 hours. Total output: 30-50 pieces.

The Variations (The 2-Hour Tuesday Refresh)

The Sunday batch produces 30-50 pieces. By Wednesday, the queue is half-empty (some pieces got more engagement, others were skipped). The Tuesday refresh adds 20-30 more pieces to the queue.

The 30-minute Tuesday refresh:

- Review what's working (top performers from Sunday's queue)

- Generate 10 variations of the top performers (AI does this in 5 minutes)

- Generate 10-20 new ideas based on what's working (AI does this in 10 minutes)

- Schedule the 20-30 new pieces

Total Tuesday time: 30 minutes. Total new content: 20-30 pieces.

The Sunday batch + Tuesday refresh = 50-80 pieces per week, 2.5-3.5 hours of creator time.

□ KEY TAKEAWAYS

- This chapter walks through the actual 2-hour Sunday batch session. The session produces 50-100 pieces of content for the week. The session has 4 phases.

- For each pillar, prompt the AI: "Give me 20 content ideas for [pillar] targeting [audience]. Each idea should be specific, not generic."

- Review the queue for balance: enough variety, not too much repetition, good distribution across pillars.

- The Sunday batch + Tuesday refresh = 50-80 pieces per week, 2.5-3.5 hours of creator time.

□ Final selection: pick 5-10 ideas for anchor content + 20-30 ideas for derivative content.

□ For each anchor piece, prompt the AI to repurpose into 5-10 platform-specific pieces:

□ Edit the draft for voice, accuracy, brand alignment (20-30 minutes per piece).

ABOUT THE AUTHOR & THE AI MASTERY FLEET

Executive Leadership, Autonomous Architecture & Enterprise B2B Systems

ENJOYED THIS FREE OPENING SAMPLE?

You have just experienced the foundational chapters of **The AI Content Agent**. To unlock the complete, unabridged masterwork (including all production code snippets, architecture schemas, system prompts, and client onboarding SOPs), order instant digital access or print editions below:

- **Official Digital Store:** <https://books.adorisedigital.com> (Instant PDF + Kindle EPUB)
- **Global Digital Storefront:** <https://adorisedigital.gumroad.com> (Worldwide Instant Access)
- **Paperback & Kindle on Amazon:** Search Google for “*The AI Content Agent Sanjay Shharma*”

Sanjay Shharma is an Author, Corporate Executive Mentor, and Technology Strategist holding an M.Sc., MBA, and APSCM from the prestigious **Indian Institute of Management (IIM) Calcutta**, with over thirty years of executive management and business leadership.

He architected **The AI Mastery & Receptionist Series** because modern enterprises are overwhelmed by repetitive operational bottlenecks, missed high-ticket phone calls, and high employee churn. Instead of theoretical AI jargon, his frameworks provide battle-tested, zero-code and low-code automated software fleets that capture revenue 24/7.

Author Credentials & Strategic Ecosystem:

Academic Alumni	Indian Institute of Management Calcutta (IIM Calcutta) • APSCM
Specialization	Autonomous AI Agents, Voice AI Receptionists & B2B Operational Scaling
Publisher	Adorise Digital • A Division of Trendy DigiStore LLC (Wyoming, USA)
Official Bookstore	https://books.adorisedigital.com
Global Storefront	https://adorisedigital.gumroad.com
Author LinkedIn	linkedin.com/in/sanjayshharma
Enterprise Portal	https://adorisedigital.com
B2B & Reader Support	support@adorisedigital.com

TURNKEY ENTERPRISE AI IMPLEMENTATION:

Need turnkey autonomous voice receptionists, lead capture bots, or programmatic marketing pipelines deployed for your firm? Explore live software architectures and book an executive consultation at adorisedigital.com or email support@adorisedigital.com.

THE 48-BOOK MASTER AUTHOR CATALOG • PART 1

Flagship Works by Sanjay Shharma • Published Worldwide by Adorise Digital

1. THE CARE COLLECTION (14 Flagship Volumes • Marital Health, Aging & Healing)

- #1 **Still Married, Just Silent** — 14 conversations to reopen an emotionally quiet marriage
- #2 **The Quiet Care** — How to give and receive care in a marriage that lasts
- #3 **The Second Chapter** — Rebuilding identity, purpose, and daily joy after caregiving
- #4 **The Second Honeymoon** — Reigniting romance, vulnerability, and play after kids leave
- #5 **The Survivor** — Compassionate protocol when one spouse is left behind
- #6 **The Long Hello** — How couples 50–70 reconnect after the last child leaves
- #7 **When They Forget Your Name** — Navigating dementia, cognitive decline, and memory care
- #8 **Gray Divorce** — Coming back together (or parting) with financial dignity at 55+
- #9 **The Last Sandwich** — Caring for aging parents, spouse, and self at the same time
- #10 **The Long Goodbye** — When your spouse is dying and the marriage is still yours
- #11 **Remarried at 60** — Blending adult families, estates, and hearts in sunset years
- #12 **After the Last Goodbye** — Practical and emotional roadmap for first 365 days of loss
- #13 **The Friends Who Fade** — Why friendships disappear in midlife & building adult tribes
- #14 **Hum Ladte Nahi (Hinglish)** — Shaadi ko wapas kholne ka 30-din ka practical plan (India)

2. THE AI MASTERY & VOICE AUTOMATION FLEET (10 Volumes • Agency & Operations)

- #15 **The AI Agent Blueprint** — Building, deploying, and monetizing autonomous software fleets
- #16 **The AI Agency Playbook** — From zero to \$50,000/month offering turnkey AI automation
- #17 **The AI Content Agent** — Autonomous video, social scheduling, and programmatic SEO
- #18 **The AI Agent Action Workbook** — Interactive architecture schemas, code snippets, and worksheets
- #19 **The AI Agent Cheat Sheets** — 50+ battle-tested system prompts for client deployment
- #20 **Never Miss Another Call** — How front-desk voice AI captures \$100,000 in lost revenue
- #21 **The AI Receptionist Blueprint** — Architecture, Twilio SIP trunking, and CRM integration
- #22 **The AI Receptionist Buyer's Guide** — What founders must know before buying or deploying
- #23 **AI Front Desk for HVAC & Home Services** — Capturing emergency dispatch calls 24/7 without missed leads
- #24 **AI Front Desk for Law Firms & Clinics** — Qualifying high-net-worth client consultations automatically

Catalog Continues on Next Page: Indian Wealth Series & Official Global Buying Links →

THE 48-BOOK MASTER AUTHOR CATALOG • PART 2

Flagship Works by Sanjay Shharma • Published Worldwide by Adorise Digital

3. THE INDIAN WEALTH & FINANCIAL MASTERY SERIES (24 Volumes • 8 Frameworks)

Books 25-27	CIBIL Score Repair Blueprint: Vol 1: 7-Step Audit Vol 2: Tactical Manual Vol 3: Legal Notices
Books 28-30	Home Loan Approval Playbook: Vol 1: Eligibility Audit Vol 2: Tactical Manual Vol 3: Banking Dossier
Books 31-33	Business Loan / MSME Finance: Vol 1: CGTMSE Audit Vol 2: Tactical Manual Vol 3: DSCR Templates
Books 34-36	Credit Card Optimization: Vol 1: APR & Fee Audit Vol 2: Tactical Manual Vol 3: Dispute Kit
Books 37-39	Debt Settlement & Relief: Vol 1: OTS Feasibility Vol 2: Tactical Manual Vol 3: Legal Defense
Books 40-42	Global Remote Career (USD): Vol 1: Resume & ATS Vol 2: Tactical Manual Vol 3: Tax & Wire Setup
Books 43-45	GST Filing & Reconciliation: Vol 1: GSTR-1/3B Audit Vol 2: Tactical Manual Vol 3: Notice Defense
Books 46-48	GST Export & Inverted Refunds: Vol 1: RFD-01 Audit Vol 2: Tactical Manual Vol 3: Appellate Appeals

WHERE TO GET THE BOOKS • OFFICIAL BUYING LINKS

1. Official Digital Bookstore (Direct DRM-Free Delivery):

- Instant access to Master PDFs, Kindle EPUBs, Audiobooks & Action Toolkits.
- Domestic Indian Payments (UPI / QR / NetBanking): <https://books.adorisedigital.com>
- Global Storefront (Cards / Apple Pay / PayPal / Whop): <https://adorisedigital.gumroad.com>

2. Amazon KDP & Google Play Books:

- Paperback, Hardcover, and Kindle editions available on Amazon worldwide.
- Simply search Google for: “[Book Title] Sanjay Shharma” or “The AI Agent Blueprint Sanjay Shharma”

3. Direct Publisher Support & Author Ecosystem:

- Email Support: support@adorisedigital.com | Portal: adorisedigital.com
- Published by **Adorise Digital** • A Division of Trendy DigiStore LLC, Sheridan, Wyoming, USA